DO NOT USE — superseded, contains a known data bug

`tpch_dataset.csv` in this directory is the **original collection output and it is wrong.** It is retained only for provenance and for the before/after comparison in `results/buffer_bug_correction_report.md`. No analysis script may read it.

The bug

`sum_buffers()` in `scripts/collection/collect_query.py` recursively summed PostgreSQL's per-node `Shared Hit/Read Blocks` up the plan tree. PostgreSQL reports those fields **cumulatively** — the root node already holds the whole-query total — so every buffer touch was counted once per ancestor level. Inflation ranges from **3.06x to 12.00x** depending on plan node count, so it cannot be undone with a single scale factor.

Fixed in the live gate by `root_buffers()` in `gate/cetp_gate.py`.

What is affected

Column	Status in this file
`shared_hit`, `shared_read`	**WRONG** — inflated 3.06x–12.00x, varies per query
`cost`, `time_ms`, `rows`, `machine_id`, `query_id`	correct, byte-identical to the corrected dataset

Because `time_ms` is unaffected, any result depending only on timings (e.g. the `naive-linear` 1/bandwidth_ratio baseline at 23.03% MAPE) is identical on both datasets. Anything touching buffers — the bottleneck classifier, `io_ratio`, `total_buffers`, and therefore the analytical roofline weights — differs.

Use this instead

`data/corrected/tpch_dataset_corrected_v2.csv` is the single canonical dataset. Every script in `scripts/analysis/` and `gate/` reads it. It is `tpch_dataset_corrected.csv` plus the `rows_processed` column, with identical values in all shared columns.

Lineage: `tpch_dataset.csv` (buggy) → `tpch_dataset_corrected.csv` (buffers repaired, see `results/buffer_bug_correction_report.md`) → `tpch_dataset_corrected_v2.csv` (+ `rows_processed`, generated by `scripts/analysis/fix1_rows_processed.py`, which is the only script that may read `tpch_dataset_corrected.csv`).